



Sequential Consistency & perl & *Perl*

About

- cPanel / WebPros
- Houston.pm.org – join us for zooms!

History

- Perl “FLAT” - Formal Language/Automata Toolkit
- Uploaded to CPAN in 2006
- Has waited since then for a real use case
- Interest in providing “concurrent” semantics in Perl that is run on a single process (perl) has revived this tool

Overview

- Perl (the language) readily admits concurrent semantics that a Human can use easily to reason – in fact, it *begs* for it
- `perl` (the interpreter/runtime) is strictly a *uniprocess* that admits **no** lightweight or shared memory concurrency
- This causes *cognitive dissonance*, which is harmful to both individuals and communities because it causes real conflict
- Reconciliation can happen if the right approach is taken
- The purpose of this talk is to demonstrate how concurrent shared memory semantics can be used for programming in a uniprocess environment (`perl` runtime)

Shared Memory Concurrency

- Communication is implicit since it depends on shared *memory* state; contrast this with explicit message passing approaches
- Allows very complex interactions among concurrent lines of execution (threads) to be described implicitly, this means less code; less code means less bugs, etc.
- Is there value in being able to write programs for `perl` that assume implicit shared memory communications; but can run “correctly” in a uniprocess? I believe, *yes*.

Motivation

“to have all the fun of traditional shared memory programming, but inside of a single process”

- **Express concurrency naturally in Perl**
- **...that can leverage the strengths of `perl`**
- **...including its implicit memory model**
- **...and `stateful` subroutines**
- **...and do it in a fun and interesting way**

Motivating Example


Using Sub::Genius

```
#  
# Implements classic JAPH :-)  
#  
  
my $pre = q{  
begin  
(  
    J &  
    A &  
    P &  
    H  
)  
end  
};
```

More on this
later ...

Important Terms to Discuss




- Uniprocess
 - Memory model
 - Sequential Consistency
 - Finite Automata
 - Regular Expression
 - Shuffle Operator
- 

perl is a *uniprocess*

- perl 5.000 (1994) had no reason to consider executing over more than a single processor
- Linux kernel 2.0 (1996) was the first step to OS support for more than 1 CPU (aka, symmetric multi-processing, SMP)
- Many attempts to *work around* this “limitation” in perl have been introduced over the years
- Most “solutions” are based on `fork`, which necessarily doesn't admit SMP
- Other solutions do *weird* things with context switching to *emulate* threads
- A LOT of **negativity** in the community as come from the lack of embracing the *uniprocess* nature of perl
- **BUT** there is a silver lining to this and I believe a path forward

...

RECAP: Perl vs perl



Sequential Consistency

With respect to “parallel” or “threaded” programs

*“... the result of any execution is the same as if the operations of all the processors were executed in **some sequential order**, and the operations of each individual processor appear in this sequence in the order specified by its program.” - Leslie Lamport [2]*

AGAIN, “all the fun and profits of traditional shared memory programming, but on fewer CPUs than intended – usually, just one (as in the case of this being applied to perl)”

Benefits of a Uniprocess

- Inherent memory consistency (no competing *threads*, no corruption)
- No race conditions
- Easy, powerful runtime memory introspection
- `perl` in particular has a *richly* developed set of variable scoping at our disposal (e.g., package, file, block, block-local, stateful subroutines)
- *Some* (me) even claim uniprocess perl has more in common with a uniprocess OS* than other PLs

*albeit one lacks a *decent* programming language; this is **not** true it has at least 6: regexes, XS, FORMAT, s/printf, un/pack, strftime

perl's *Memory* Model

- In computing, a memory model describes: the interactions of threads through memory and their shared use of the data. (wikipedia)
- perl is a uniprocess; there are no threads; it is sequential – so what's the “memory model”?
- The closest we have are Perl's *variable scoping*, data isolation, refcounting, etc
- perl also can give subroutines memory between calls via the `state` declaration

What about Performance

- This talk doesn't consider this directly
- And when most people discuss “performance” in the context of “concurrency”, what they really care about is “*speed up*”
- Speed up; i.e., the decreasing “*time to solution*” one achieves as more processes are thrown at a particular task – see `fork`
- We shall see that the approach presented offers some interesting opportunities for managing “speed up”

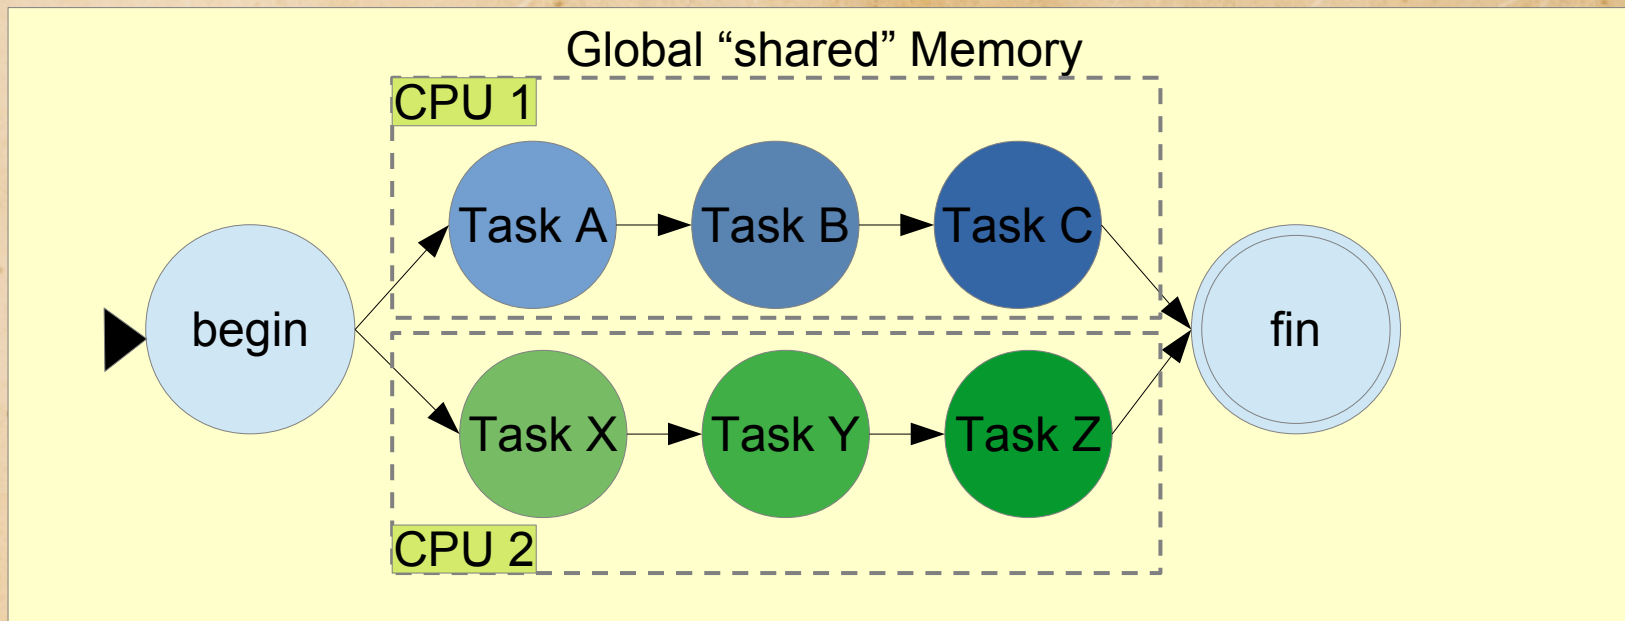
Where this is heading: Concurrent SMP

- Shared memory accomplished through global variables
- Race conditions reduce down to “*safe*” side effects
- No competition for system resources (e.g., disk i/o, network, etc)
- No ‘corrupted’ memory due to partial writes by 2 or more “threads”
- Though, “overwrites” are absolutely possible; the writes will still be complete
- All the joins of SMP programming and *none* of the speed up! - So what’s the benefit?

Sequential Consistency

Example by dependency graph, "A→B→C & X→Y→Z"

- Imagine a program that is written (in a supportive language) such that two independent lines of execution exist.



Sequential Consistency

Example by dependency graph, " $A \rightarrow B \rightarrow C$ & $X \rightarrow Y \rightarrow Z$ "

- *Total Ordering?*

A must precede B, B must precede C


X must precede Y, Y must precede Z

- *Partial Ordering?*

Tasks in $\{A, B, C\}$ may happen in any order with respect $\{X, Y, Z\}$

“execution” paths $A \rightarrow B \rightarrow C$ is independent of the path $X \rightarrow Y \rightarrow Z$


This independence *implies* concurrency



Sequential Consistency

Example by dependency graph, "A→B→C & X→Y→Z"

- “ABC” and “XYZ” naturally map to an environment with, e.g., 2 CPUs
- Each independent execution path may happen simultaneously
- Using more than 2 CPU units provides no parallel “*speed up*”




- What about using 1 processing unit?

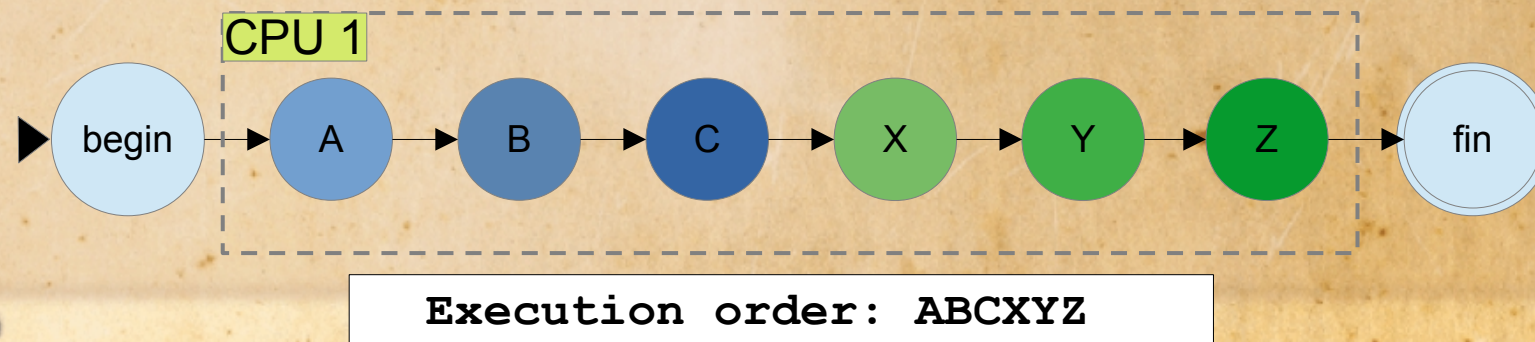
How Sequentialize Consistently?

Example by dependency graph, " $A \rightarrow B \rightarrow C$ & $X \rightarrow Y \rightarrow Z$ "

- How do we go from 2 CPUs:



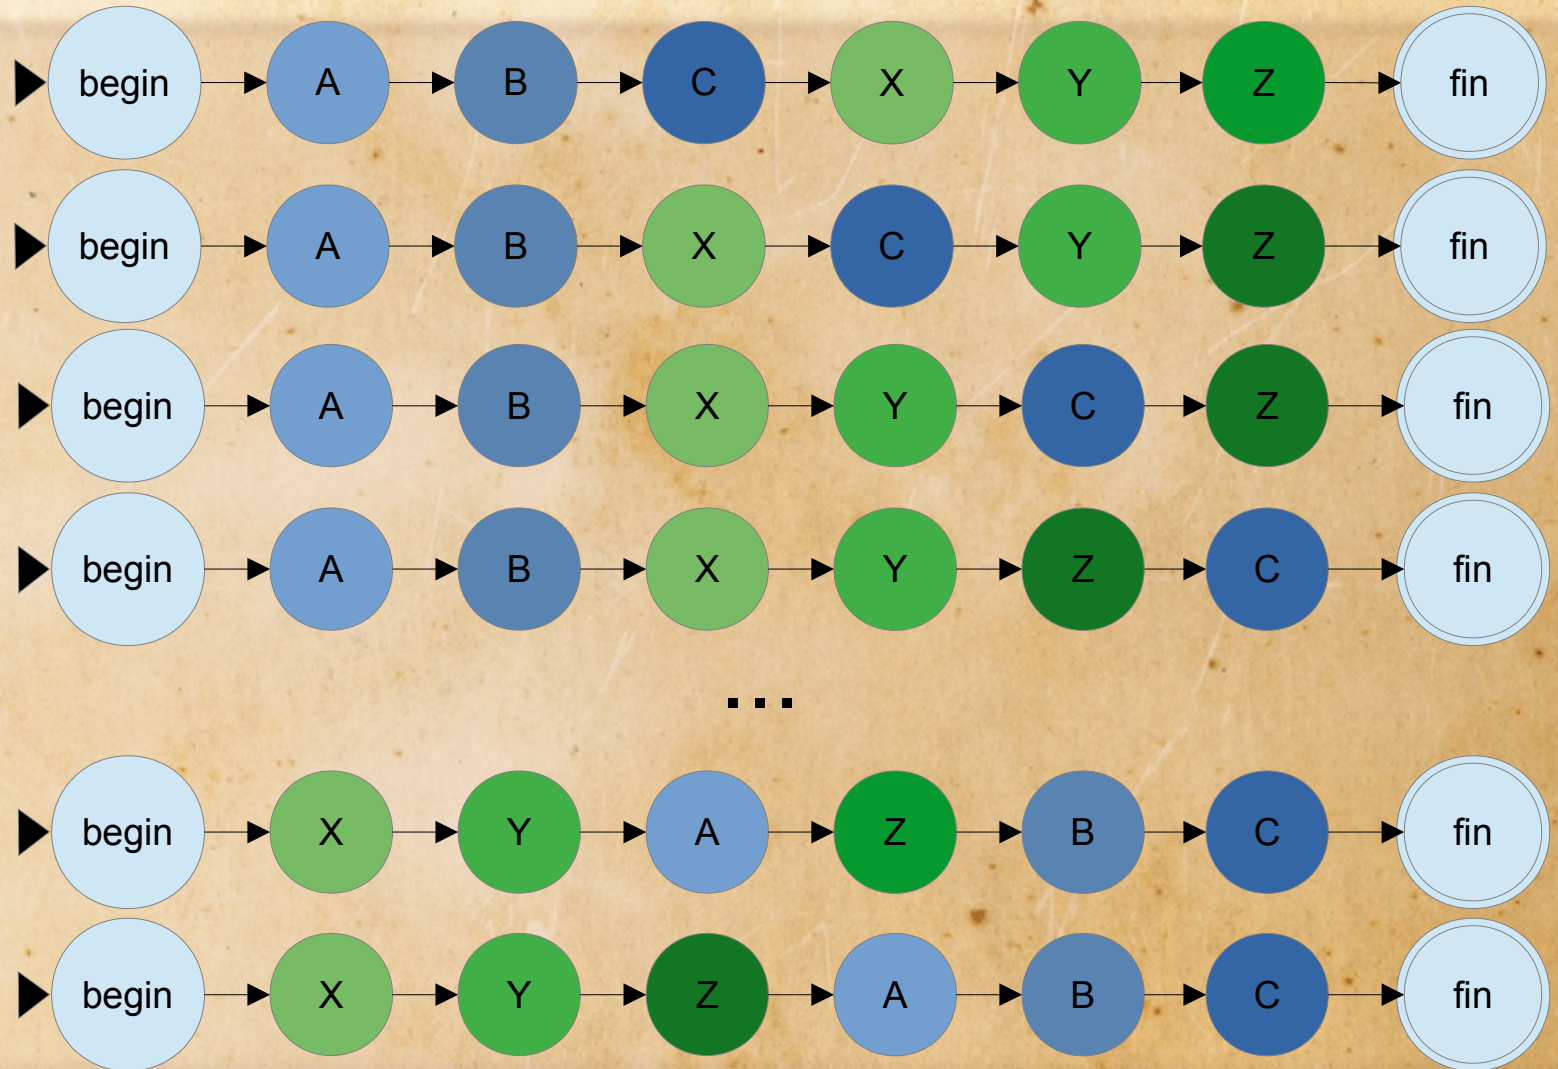
- To 1 CPU such that the serialized execution is consistent with a valid ordering using 2 CPUs?



Sequential Consistency

Example by dependency graph, "A→B→C & X→Y→Z"

ABCXYZ
ABXCYZ
ABXYCZ
ABXYZC
AXBCYZ
AXBYCZ
AXBYZC
AXYBCZ
AXYBZC
AXYZBC
XABCYZ
XABYCZ
XABYZC
XAYBCZ
XAYBZC
XAYZBC
XYABYZ
XYABZC
XYAZBC
XYZABC




However, there are **20** valid sequential orderings implied by the concurrent "version" executing in parallel!

Sequential Consistency

Example by dependency graph, "A→B→C & X→Y→Z"


ABCXYZ
ABXCYZ
ABXYCZ
ABXYZC
AXBCYZ
AXBYCZ
AXBYZC
AXYBCZ
AXYBZC
AXYZBC
XABCYZ
XABYCZ
XABYZC
XAYBCZ
XAYBZC
XAYZBC
XYABCZ
XYABZC
XYAZBC
XYZABC



- All *sequentially consistent* (i.e., valid) execution orderings maintain the total orderings: "A→B→C" & "X→Y→Z"
- Yet, admit the partial orderings permitted by the lack of any dependency between {A,B,C} and {X,Y,Z}.

How Can We Serialize *Arbitrary Dependencies*?

- **Step 1** - recognize a dependency graph is just a directed *graph*



- **Step 2** – recognize that the serialization of a graph consisting of nodes and labels has been studied a lot over the decades
- **Step 3** – recall that day in your CS class on Computer Theory they covered the *shuffle* operator

Spoiler! We can do so using well known concepts and algorithms from the study of Finite Automata and Regular Language.

Finite Automata

- **Deterministic Finite Automata (DFA)**

A state machine that can be described as a regular expression (not the Perl kind!)

Describes a set of valid strings; also known as a *Regular Language*

- **Non-deterministic Finite Automata (NFA)**

A DFA + an ε -transition (“epsilon”) that admits non-determinism

- **Parallel Finite Automata (PFA)**

A NFA + a λ -transition that adds additional constraints to the language it describes

Equivalent to binary PetriNets (*Stotts, Pugh 1994 – Parallel Regular Expressions*)

Describing Regular Languages

- **Deterministic Finite Automata (DFA)**

Concatenation, Union, Kleene Star

- **Non-deterministic Finite Automata (NFA)**

Concatenation, Union, Kleene Star

- **Parallel Finite Automata (PFA)**

Concatenation, Union, Kleene Star

+ the *Shuffle* operator

A PFA is equivalent to a *binary* PetriNet

What is the Shuffle *Operator* [1]?

- Most are familiar with operations over regular languages:

Concatenation (*usually* implied, “.”)

Union (“|”)

Kleene Star (“*”)

- Under these operations, the results are “closed” under RLs; i.e., their application results in another RE
- Shuffle is a less well know, but properly closed under REs.
- Is the “and” to Union’s “or”
- Effectively *interleaves* valid strings from two regular languages
- A RE with a shuffle is called a, *Parallel Regular Expression*
- *Operator* is often represented as an ampersand, “&”

What is the Shuffle Operator?

Examples

- Language 1 (L_1) – the string “**ab**”
- Language 2 (L_2) – the string “**cd**”
- Shuffle, L_1 & L_2 :

“**ab&cd**”

Requires a valid string from BOTH languages are *interleaved* in any valid string

Valid strings are:

abcd

acbd

acdb

cadb

cdab

Converting Non-Shuffled REs to DFA

- RE \rightarrow NFA

Recursive descent parse to create an AST

Thompson Construction to construct a NFA from a traversal of the AST


Introduces ϵ -transitions

- NFA \rightarrow DFA

Subset Construction Algorithm

- DFA \rightarrow valid strings

Depth-first graph traversal



Converting Shuffled PREs to DFA

- PRE \rightarrow PFA

Recursive descent parse to create an AST

“*modified*” Thompson Construction to construct a PFA from a traversal of the AST

Introduces both ϵ -transitions and λ -transitions

- PFA \rightarrow NFA

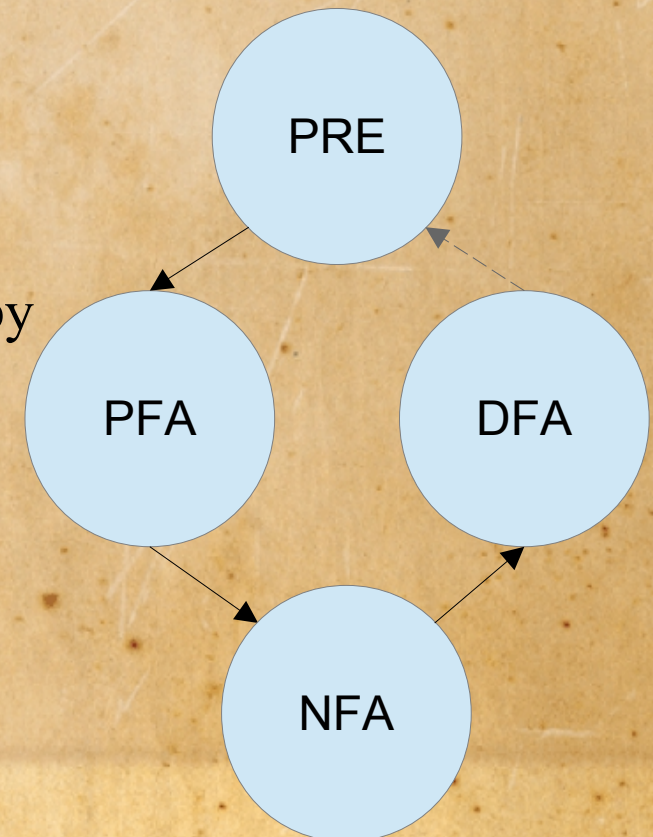
Algo replaces λ -transitions with ϵ -transitions by enumerating valid partial orderings

- NFA \rightarrow DFA

Subset Construction Algorithm

- DFA \rightarrow valid strings

Depth-first graph traversal



Strings from a *finite* Language

- A *finite* language is any language described by a Regular Expression that does **NOT** include a **Kleene Star**; otherwise there are an *infinite* set of valid strings
- *Given* a DFA, generate a set of valid strings in the language it describes
- Solution:

traverse graph from start to an accept node, collect labels (e.g., letters or symbols) along the way

Print set of collect symbols collected along the path


Strings from an *infinite* Language

- *(slide added for completeness, no examples in this talk use infinite FAs to produce execution schedules)*
- An *infinite* language is any language described by a Regular Expression that **does** include a **Kleene Star**; otherwise there are an *finite* set of valid strings
- *Given* a DFA, generate a set of valid strings in the language it describes
- Solution:
 - traverse graph from start to an accept node, collect labels (e.g., letters or symbols) along the way
 - Manage how “deep” into an infinite cycle the traversal is, back out once a maximum is met
 - Print set of collect symbols collected along the path

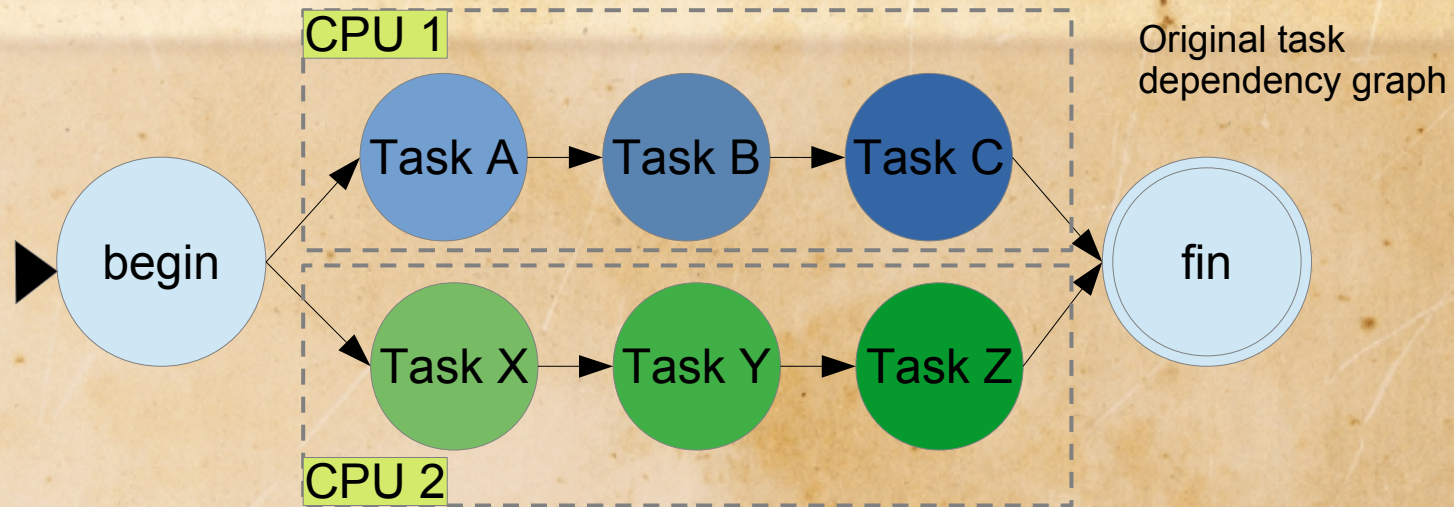
Shuffle Example

Example of “parallel” Regular Expression, “ABC&XYZ”


- “ABC” & “XYZ”
- More formal (labeled transitions, numbered states) diagram of the **Parallel Finite Automate**:



Resulting PFA




**PFA for
(ABC) & (XYZ)**

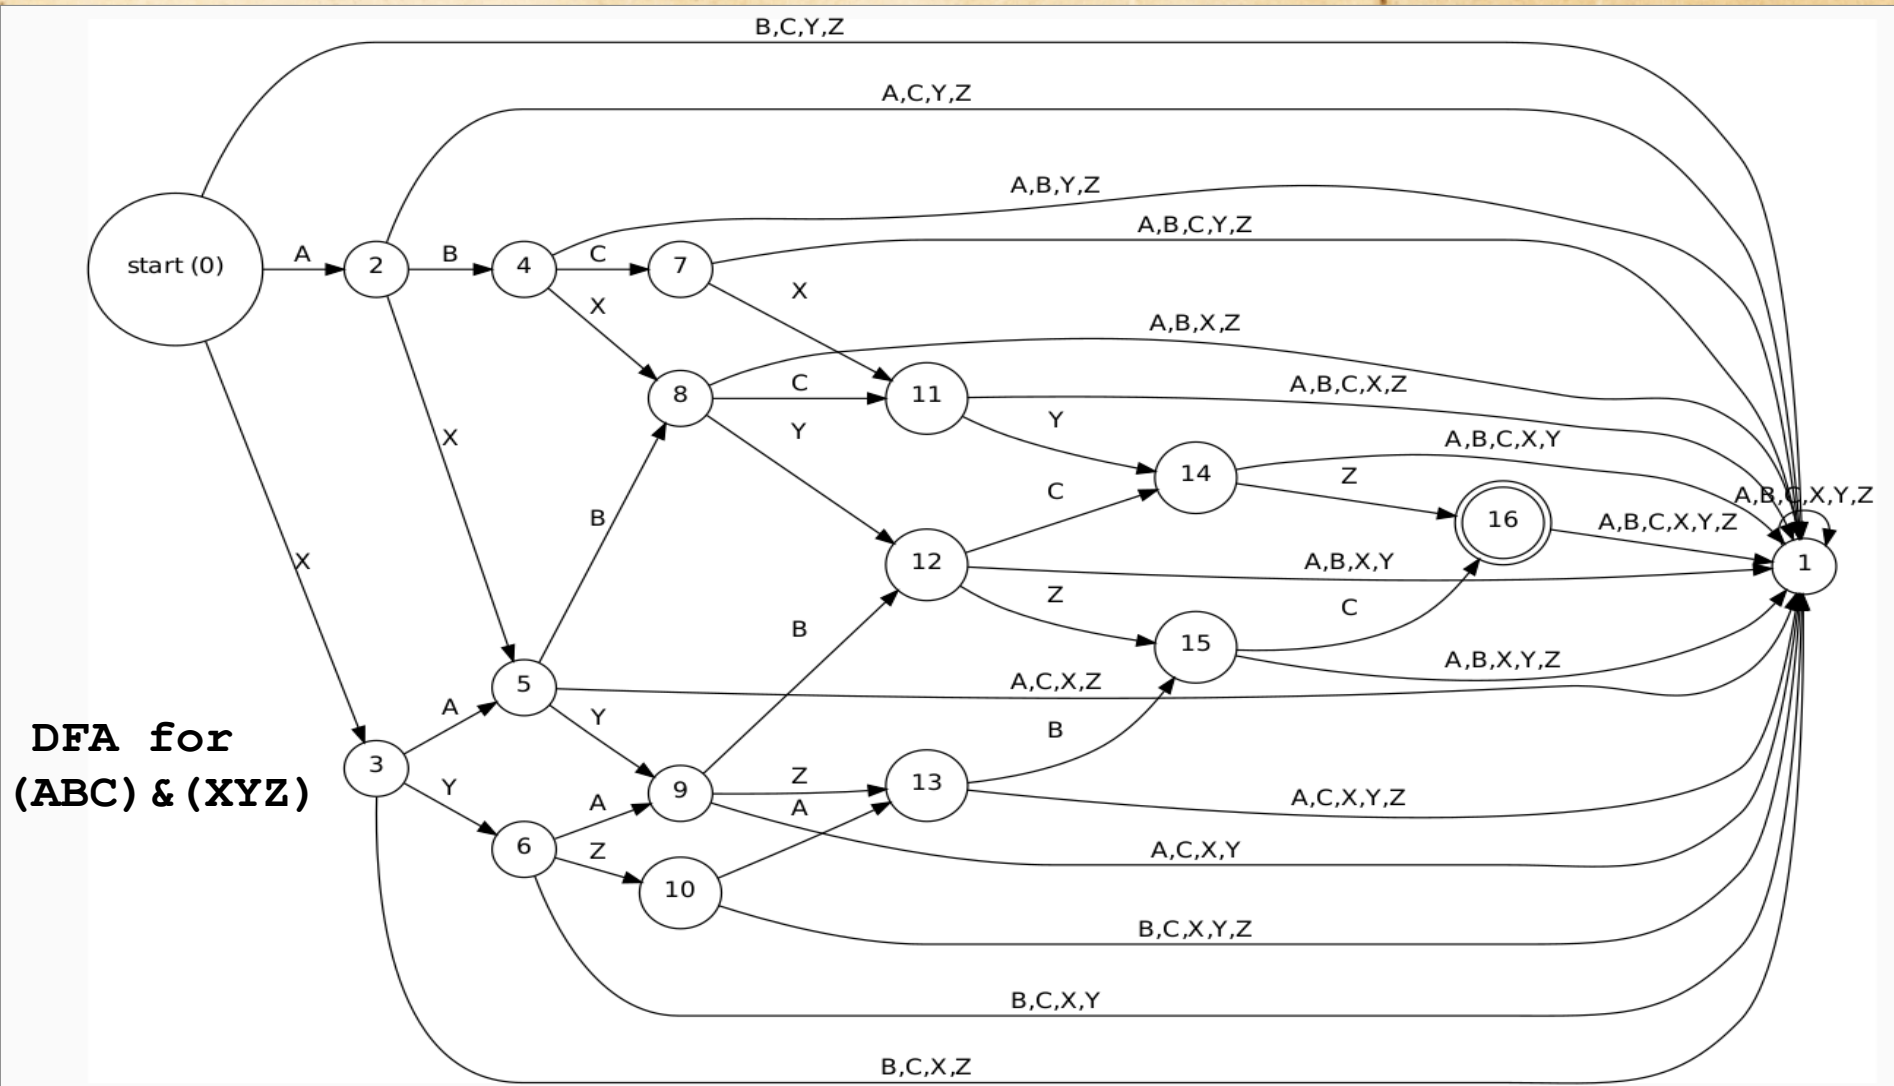


Resulting NFA

**NFA for
(ABC) & (XYZ)**



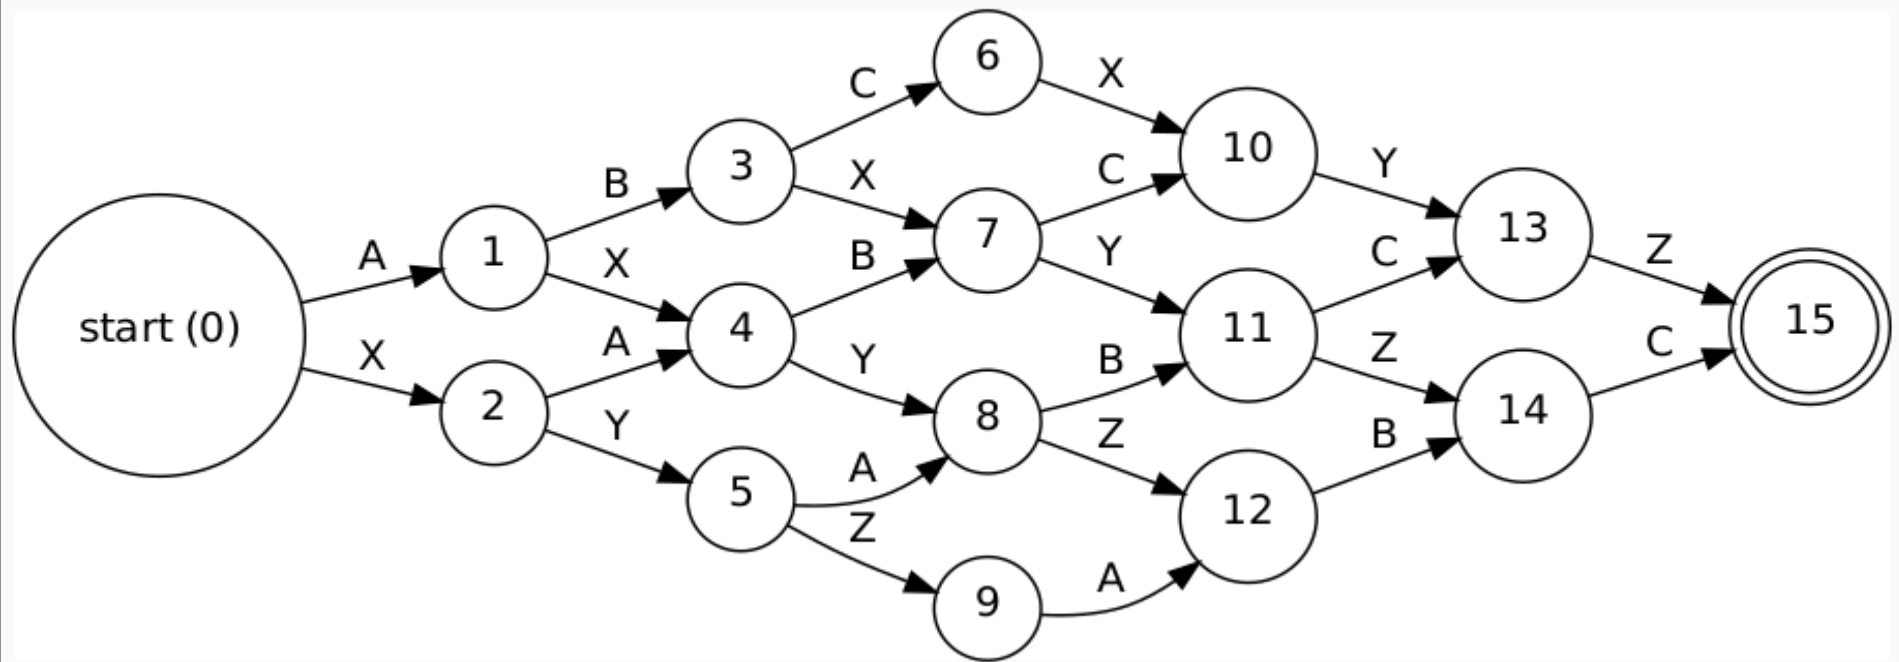
Resulting* DFA



*after subset construction NFA→DFA



Resulting Minimized DFA

min DFA for
(ABC) & (XYZ)



Enumeration of all Start → Accept Paths

min DFA for
(ABC) & (XYZ)




ABCXYZ
ABXCYZ
ABXYCZ
ABXYZC
AXBCYZ
AXBYCZ
AXBYZC
AXYBCZ
AXYBZC
AXYZBC
XABCYZ
XABYCZ
XABYZC
XAYBCZ
XAYBZC
XAYZBC
XYABYZ
XYABZC
XYAZBC
XYZABC

All Start → Accept paths represent valid strings resulting from the interleaving of valid strings from the *shuffled* “languages”.


Because this process maintains the partial ordering required by each language in the shuffle, these strings are sequentially consistent, and therefore represent valid execution orders implied by the original concurrent description.

Example with all *Finite* RE Operators

PFA for
 $A(B|C) \& X(Y|Z)$



NFA for
 $A(B|C) \& X(Y|Z)$



Recall: *Motivation*

- **Express concurrency naturally in Perl**
(for Humans to reason easily)
- **...that can leverage the strengths of perl**
(uniprocess, memory models)
- **...and do it in a fun and interesting way**
(primary goal)

Sub : : Genius

A Perl module for managing concurrent semantics for serial execution

- Takes a concurrent “plan”, i.e., a Parallel Regular Expression
- “dumbs down” SMP for perl, hence the module name (credit: *TEODESIAN*)
- A “plan” describes the execution dependency of subroutines with concurrent semantics (i.e., using a “*shuffle*” operator)
- Serializes it to one of likely *many* serialized forms
- All subroutines are called in the same `perl` process, so we can leverage:

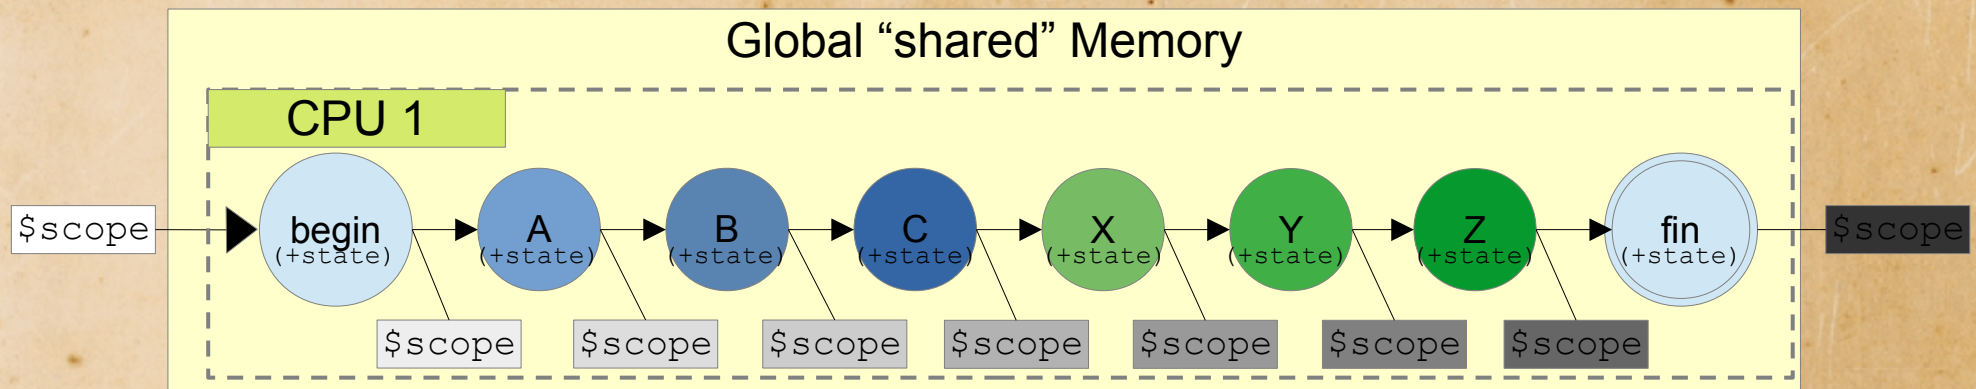
Global (*shared*) memory

Execution “scope” passed in/out of each subroutine call

Subroutines with **state** that may change from call to call!

Sub : : Genius

A Perl module for managing concurrent semantics for serial execution




- Share memory space, accessible to all subroutines
- An execution "\$scope" that is passed in a *pipeline* fashion
- Subroutines that "remember" between calls – e.g., coroutines

Anatomy of a Sub::Genius program

minimal DFA for

begin (step1 & step2) fin



Valid sequential forms:

- begin → step1 → step2 → fin
- begin → step2 → step1 → fin

subroutines with “state”
adds co-routines

subroutine implementations

```
use strict;  
use warnings;  
use feature 'state';
```

```
use Sub::Genius ();
```

```
# the plan is a PRE, 'parallel regular expression' in the Formal sense  
my $preplan = q{  
  begin  
    ( step1 & step2 )  
  fin  
};
```

declare “parallel plan”

```
# Load PRE describing concurrent semantics  
my $sq = Sub::Genius->new( preplan => $preplan );
```

```
# treated as “shared memory”
```

```
my $GLOBAL = {};
```

shared memory

```
# run plan
```

```
my $final_scope = $sq->run_any( scope => {}, );
```

```
# implement sub routines
```

```
sub begin {  
  my $scope = shift;  
  state $sub_state = {};  
  # ..  
  return $scope;  
}
```

```
sub step1 {  
  my $scope = shift;  
  state $sub_state = {};  
  # ...  
  return $scope;  
}
```

```
sub step2 {  
  my $scope = shift;  
  state $sub_state = {};  
  # ...  
  return $scope;  
}
```

```
sub fin {  
  my $scope = shift;  
  state $sub_state = {};  
  # ...  
  return $scope;  
}
```

execution “scope”

Sub :: Genius Tooling & Info



- stubby

General utility for making it easier to create and work with programs utilizing the sequential planning of Sub :: Genius

- fash


Wrapper around FLAT one-liners; old and crusty, but still useful for learning things about PREs and the FA they define (which can be surprising)

- S :: G has a lot of information in the POD (perldoc Sub :: Genius), but more is being written

Sub : : Genius's Big Wart

- PREs take exponentially long to convert into a DFA (algo complexity is largely unavoidable)
- Solution:

```
stubby precache -p "a&b&c&d&f"
```
- Basically a “compile” step; saves using Storable.
- This programming model for Perl/perl is still embryotic, but it's interesting and powerful
- More work must be done to explore this space



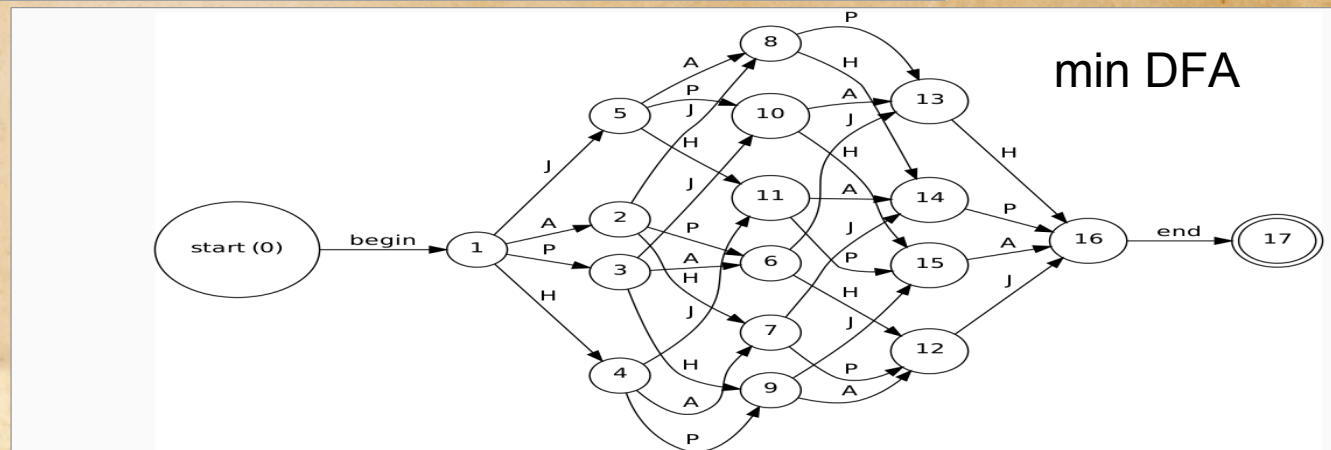
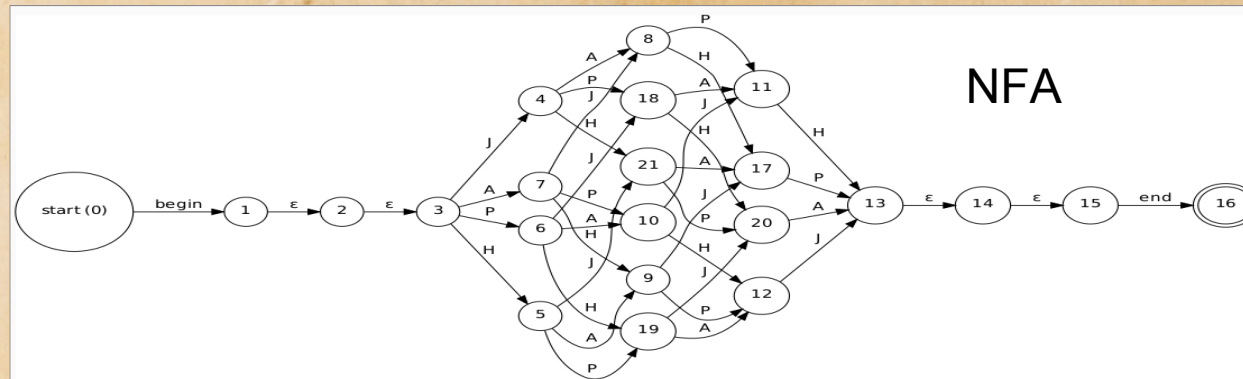
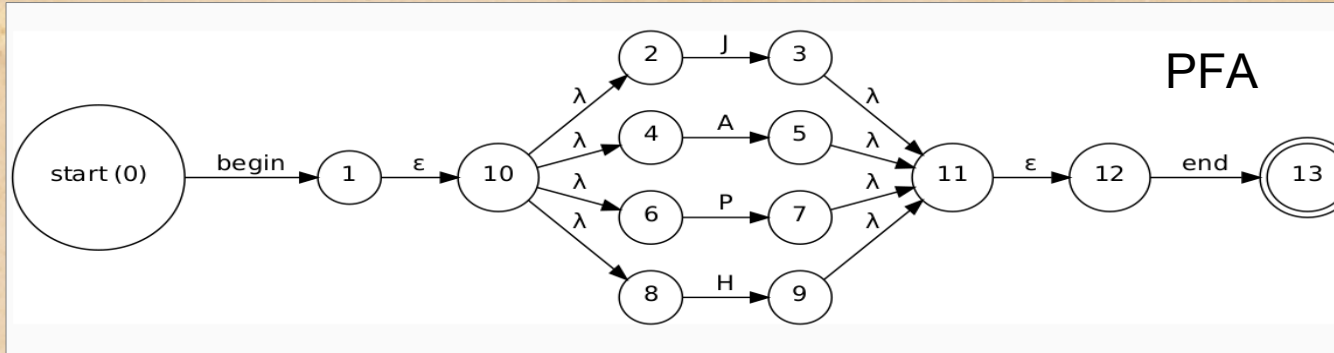
**"Now is the time on
Sprockets when we
demo."**

just another Demo

```
#  
# Implements classic JAPH :-)  
#  
  
my $pre = q{  
begin  
(  
    J &  
    A &  
    P &  
    H  
)  
end  
};z
```


Underlying Automata

Example – “JAPH”, stages of serialization using FA conversion algorithms



A Way Forward for...


- Allows for use of concurrent semantics *now*
- Exploitation of `perl`'s uniprocess memory model
- Exploring concurrent semantics in Perl – e.g., “fork/join”, “async/away”, etc
- Other things, I am sure

SEE ALSO

- Pipeworks
- Sub::Pipeline
- Process::Pipeline
- FLAT
- Graph::Petrinet


Good Reads




- 1. <https://www.planetmath.org/shuffleoflanguages>
 - 2. Leslie Lamport, "How to Make a Multiprocessor Computer That Correctly Executes Multiprocess Programs", IEEE Trans. Comput. C-28,9 (Sept. 1979), 690-691.
 - 3. <https://www.hpl.hp.com/techreports/Compaq-DEC/WRL-95-7.pdf>
 - 4. <https://troglodyne.net/video/1615853053>
 - 5. Introduction to Automata Theory, Languages, and Computation; Hopcroft, Motwani, Ullman. Any year.
- 

Conclusion



- PREs can be used to describe on a useful level, the execution plan of subroutines for the purpose of auto-serialization
 - Useful for expressing SMP concurrency for valid executions on single processes
 - Lots of work still needs to be done to explore the benefits
- 

Open Questions



- What does it look like when fork is introduced to the supporting subroutines?
 - What what is possible when subroutines are treated statefully? (e.g., subs implicitly retain memory)
 - How does “async/await”, “futures”, etc translate using this approach
 - How can “infinite” languages be used effectively
 - What runloop/driver controls would be useful?
 - What utilities (e.g., stubby, fash) would be useful for development and verification of programs written using this approach
- 